

The Defaults Ledger

SUBThe Defaults Ledger

p-rick research program

p-rick — an independent research program on missing software: Disclosure and governance for population-scale configuration.

The Defaults Ledger

SUBThe Defaults Ledger

p-rick research program

Working Paper P-016 · v1.0 · October 7, 2026 · CC BY 4.0

Abstract — The default is the most powerful user interface element ever built, and it governs with no records, no disclosure, and no ledger. The behavioral evidence is a half-century deep and uncomfortably large: organ-donation consent collapses from ~86–99% to single digits when a country moves default from opt-out to opt-in (Johnson & Goldstein, 2003, the canonical demonstration); retirement-plan enrollment jumped from under 50% to 86% when enrollment became the default (Madrian & Shea, 2001); encryption for a billion users arrived not through persuasion but through one company making it the default (WhatsApp, 2016); a single new default-on prompt reallocated an estimated \$10 billion of advertising economics in a year (Apple’s App Tracking Transparency, 2021, with Meta’s own published estimate of the impact); and the market’s price for a default is now public — Google’s payments to be the default search engine on Apple devices reached roughly \$20 billion in 2022, disclosed in antitrust exhibits. Against this measured power, the practice of setting defaults is an unrecorded act: platform teams change defaults silently in ordinary releases, users discover the changes as behavior drift (“why is my phone doing this now?”), researchers must reverse-engineer default changes after the fact, regulators legislate individual settings one at a time, and the enterprise world — which solved default governance decades ago for exactly this reason (Group Policy, MDM configuration profiles, managed preferences) — has never ported the machinery to civilians. The result is a power asymmetry with no instrument: defaults are legislation nobody voted on, executed as configuration, invisible as policy. We specify the **defaults ledger**: machine-readable **defaults manifests** per application version (setting, value, source, rationale, changed-from, changed-when), with **default provenance** (who set it, under whose decision, with what stated reason — the commit history for policy); **diff and drift notification** (version-to-version default changes surfaced to the user at update time in the same notification surface that announces the update); a **personal override registry** (your deviations from every product’s defaults, portable, exportable, reinstall-surviving — the setting that follows you across devices and upgrades); **negotiation layers** (household and organizational default policies, importable default packs from trusted third parties — the editorial layer the standards never had); **default impact assessments** for changes above an exposure threshold (the environmental-impact-statement pattern applied to configuration: expected behavioral displacement, affected population, reversal path); and **procurement-grade defaults bills of materials** (the artifact a CIO, a regulator, or a browser-vendor reviewer can diff). Every component exists in an adjacent world — enterprise config management, release-note practice, package-lock manifests, regulatory impact assessment, the web-standards governance culture — and no platform, standard, or product composes them for the settings that govern billions of non-enterprise users. The gap is STRONG: the effect sizes are the largest in behavioral science, the asset (default placement) is actively bought and sold at disclosed billions, the governance instruments are proven in the enterprise, and the civilian layer is a wilderness of invisible unilateral amendments. Defaults are the legislation of the default polity. This paper specifies the parliamentary record.

Keywords: p-rick

p-rick working paper P-016 · series IV (assumptions) · draft 1.0

The default is the most powerful user interface element ever built, and it governs with no records, no disclosure, and no ledger. The behavioral evidence is a half-century deep and uncomfortably large: organ-donation consent collapses from ~86–99% to single digits when a country moves default from opt-out to opt-in (Johnson & Goldstein, 2003, the canonical demonstration); retirement-plan enrollment jumped from under 50% to 86% when enrollment became the default (Madrian & Shea, 2001); encryption for a billion users arrived not through persuasion but

through one company making it the default (WhatsApp, 2016); a single new default-on prompt reallocated an estimated \$10 billion of advertising economics in a year (Apple’s App Tracking Transparency, 2021, with Meta’s own published estimate of the impact); and the market’s price for a default is now public — Google’s payments to be the default search engine on Apple devices reached roughly \$20 billion in 2022, disclosed in antitrust exhibits. Against this measured power, the practice of setting defaults is an unrecorded act: platform teams change defaults silently in ordinary releases, users discover the changes as behavior drift (“why is my phone doing this now?”), researchers must reverse-engineer default changes after the fact, regulators legislate individual settings one at a time, and the enterprise world — which solved default governance decades ago for exactly this reason (Group Policy, MDM configuration profiles, managed preferences) — has never ported the machinery to civilians. The result is a power asymmetry with no instrument: defaults are legislation nobody voted on, executed as configuration, invisible as policy. We specify the **defaults ledger**: machine-readable **defaults manifests** per application version (setting, value, source, rationale, changed-from, changed-when), with **default provenance** (who set it, under whose decision, with what stated reason — the commit history for policy); **diff and drift notification** (version-to-version default changes surfaced to the user at update time in the same notification surface that announces the update); a **personal override registry** (your deviations from every product’s defaults, portable, exportable, reinstall-surviving — the setting that follows you across devices and upgrades); **negotiation layers** (household and organizational default policies, importable default packs from trusted third parties — the editorial layer the standards never had); **default impact assessments** for changes above an exposure threshold (the environmental-impact-statement pattern applied to configuration: expected behavioral displacement, affected population, reversal path); and **procurement-grade defaults bills of materials** (the artifact a CIO, a regulator, or a browser-vendor reviewer can diff). Every component exists in an adjacent world — enterprise config management, release-note practice, package-lock manifests, regulatory impact assessment, the web-standards governance culture — and no platform, standard, or product composes them for the settings that govern billions of non-enterprise users. The gap is STRONG: the effect sizes are the largest in behavioral science, the asset (default placement) is actively bought and sold at disclosed billions, the governance instruments are proven in the enterprise, and the civilian layer is a wilderness of invisible unilateral amendments. Defaults are the legislation of the default polity. This paper specifies the parliamentary record.

1 Introduction

In 2003, Eric Johnson and Daniel Goldstein published a one-page paper in *Science* that has quietly organized twenty years of policy since: “Do Defaults Save Lives?” Across European countries, effective organ-donation consent rates ran from 4.3% to 27.5% in opt-in nations — and from 85.9% to 99.98% in opt-out nations. Austria and Germany, culturally adjacent neighbors with near-identical populations and attitudes toward donation, differed by a factor of eight in consent — and the entire difference was the default. The paper’s abstract is one sentence long, and its content is the governing fact of this paper: defaults are not suggestions; they are *outcomes*, manufactured at population scale, by a line of configuration.

Software runs on the same physics with the same magnitudes. Retirement enrollment at 86% versus 49%: the default. A billion conversations encrypted in a day: one default flip, 2016, and every subsequent negotiation about the encryption of ordinary messaging happened on terrain that flip had created. A trillion-dollar advertising apparatus repriced in a year: one default-on prompt, 2021. The right to place a search box’s default: worth roughly \$20 billion a year — not estimated by academics but disclosed in antitrust exhibits, a public market clearing price for the single most consequential bit of configuration in the consumer economy. The default is where a company’s entire preference architecture becomes the population’s behavior, without a single user decision being made.

And the act of setting one is recorded nowhere.

A platform flips a default in an ordinary release. The release notes — if the change appears at all — carry it as a feature (“Now more private!” “Now more personalized!”). The user’s behavior changes without their knowledge; their battery life, their privacy surface, their content exposure, their auto-upload topology silently reconfigures. Six months later, a researcher or a journalist or a congressional staffer reconstructs the change after the fact, if anyone ever does. Regulators respond by legislating *the individual default* (this one shall be off; that one shall be on) — a setting-by-setting whack-a-mole that requires the full political apparatus of a nation per bit, while the *practice* of unilateral silent default amendment — the machinery that produces thousands of such bits per year across the ecosystem — is never itself governed. The enterprise world, which long ago learned that configuration is power, built itself an entire discipline of default governance: Group Policy, MDM profiles, managed preferences, config baselines, change review. The civilian — the user whose phone holds two thousand settings they have never seen — has exactly one governance instrument: “reset to defaults,” which is not governance but surrender.

The thesis in one sentence: **defaults are population-scale legislation executed as configuration, and the governance they require — manifests, provenance, diffing, drift notification, override registries, negotiation, and impact assessment — is a fully specifiable software layer that already exists in enterprise fragments and has never been composed for the public.**

2 A theory of default power

Why is one bit of configuration worth \$20 billion a year? The power is a compound of four mechanisms, each independently documented in the behavioral literature, and each individually sufficient to make the default the highest-leverage element in any interface:

1. **The pre-decision.** A default is a decision made on the user’s behalf before the choice point arrives. The user who never visits the setting page has delegated the decision completely; the delegation is invisible because it is silence-shaped. Status-quo bias and the endowment of the pre-chosen make reversal psychologically expensive relative to the initial non-choice.
2. **Friction asymmetry.** Accepting the default costs zero actions; departing costs attention, navigation, comprehension, and risk (what else does this toggle touch?). The behavioral result is that option architecture beats option preference wherever the friction gap is large — which is the entire finding of the organ-donation and retirement-enrollment literatures at population scale.
3. **The endorsement channel.** Users read defaults as *implicit expert recommendations* — “the company must think this is the right setting” — a rational heuristic under time scarcity that converts configuration into authority. A default is a statement the company never has to sign.
4. **The network compounding.** Many defaults coordinate behavior (messaging protocols, sharing visibility, content exposure, app-install permissions), so the default *is* the standard: whoever sets it sets the topology of the ecosystem that forms around it — the encryption-by-default case being the cleanest example, where one company’s default became the industry’s baseline security posture.

The four mechanisms together explain the market price. What they do not explain — because nothing in the machinery explains it — is why the act of wielding this power is unrecorded, undisclosed, and unreviewed. Power with these effect sizes and zero accounting is not a design choice; it is a governance vacuum.

3 The landscape: solved in the enterprise, wilderness in the public

Enterprise configuration governance. The complete pattern library, mature for two decades: directory-backed policy (Group Policy since 1999), mobile-device management profiles, configuration baselines (CIS benchmarks), change-management review for policy changes, and audit trails. The enterprise solved default governance because enterprises *pay* for ungoverned configuration — in breaches, in compliance findings, in coordination costs. Every instrument this paper needs for civilians exists here, priced for organizations, never ported down.

Package management and lockfiles. The software supply chain built its own defaults ledger without naming it: the lockfile is a recorded, diffable statement of *what was selected by default* (resolved, not chosen) with provenance to versions and hashes; `diff` between lockfiles is the drift notification; the manifest format is standardized and machine-readable. The ecosystem around code treats dependency defaults as *records to be governed*; the ecosystem around user settings treats them as internal implementation detail.

Release-note and changelog practice. The disclosure culture exists in embryo: disciplined projects log behavior changes; browsers maintain enterprise-readable release notes specifically because *enterprises demanded* default-change visibility, which is the demand this paper generalizes to everyone. The consumer-facing surface, meanwhile, markets the change (“new! improved!”) rather than disclosing it (“your default changed”).

Regulation of individual defaults. The legislative pattern: GDPR legislating consent defaults; specific mandates for specific bits; the U.S. state privacy statutes enumerating required-off defaults category by category; the EU’s interop and browser-choice remedies (the browser-ballot saga — a regulator literally manufacturing a default-choice moment as an antitrust remedy, twice, with mixed results that prove how hard it is to *legislate choice architecture post hoc* when the platform owns the layer it executes on). The pattern’s limit is the setting-by-setting granularity: each bit requires its own political event, while the practice — silent unilateral amendment at release cadence — continues underneath.

Default research. A rich behavioral literature (the Johnson-Goldstein lineage; defaults in privacy decisions; default effects in pro-social and health contexts) that measures magnitudes, documents the endorsement heuristic, and debates mechanism — all at the *effect* level. The governance of the *act* is not in the research portfolio, because no instrument existed to study it with; there is no dataset of default changes to analyze, because nobody records them.

Privacy and tracking communities. The folk archivists: researchers who monitor tracking surfaces, security researchers who diff app binaries for new default-on telemetry, journalists who notice a billion phones changed behavior. The existence of the practice — reverse-engineering the ledger from outside — is the strongest evidence the ledger does not exist.

The platform’s own internals. Every platform already maintains, internally, the machinery a defaults ledger needs: feature-flag systems with owners and rationales, config databases (one OS literally names its settings store `defaults`), A/B records that measure default-change impacts with scientific precision *before* shipping and discard them after. The information exists; it is the disclosure that does not. macOS’s `defaults` command — the reading of a settings store by its platform’s own name — is the landscape’s best joke: a defaults database with no defaults governance.

VERDICT: STRONG. The power is measured at the largest effect sizes in behavioral science; the asset is priced in public markets; the governance machinery is proven in the enterprise and the supply chain; the civilian layer is a vacuum; and the disclosure practice exists in enough embryonic forms (release notes, lockfiles, internal flag ownership) that the composition is a decision, not an invention. No incumbent ships a defaults manifest for consumers. The gap is the paper.

4 Why the gap survived

The default is the product’s leverage, and disclosure prices the leverage away. The commercial value of a default runs through its invisibility: the endorsed share of traffic is worth billions *because* the user never faces the choice at full friction. A manifest, a diff, a drift notification — each converts silent amendment into a visible decision the user may actually make, and the expected value of visibility to the *setter* is negative for every default whose value derives from non-choice. The parties with the ledger’s building blocks are the parties the ledger would govern. This is the cleanest incentive inversion in this program’s catalogue, and it predicts the ledger will arrive either from regulation, from a competitive honesty market (a browser, a platform, a vendor deciding disclosure is a feature), or from the research community forcing it into the light — the three arrival channels this program’s papers have learned to expect for exactly this gap shape.

Civilian demand is diffuse and unaggregated. The user’s experience of a silent default change is “my phone got weird after the update” — an attribution failure that never organizes into political demand, because the user cannot see the bit that moved. You cannot demand the ledger for a change you cannot name. The gap is self-concealing at the individual level, the same concealment property as P-015’s estimation bias: the harm exists as ambient discomfort, not as an attributable event, and unattributable events do not become constituencies.

The enterprise precedent has a boundary, and everyone respects it. Group Policy governs defaults for *managed* devices, where a principal (the employer) owns the configuration relationship and demanded governance from vendors. The civilian owns their configuration relationship too — but alone, at $n=1$, without procurement leverage, never organized into the demand pattern vendors respond to. The machinery was built behind the corporate boundary and never exported, because the demand signal behind the boundary was collective and the signal outside it is not.

Regulators inherited the setting-by-setting culture. Consumer-protection law developed per-harm: this practice is unlawful, that disclosure is required, the other toggle is mandatory. The practice-level instrument — govern the *ledger of amendments* rather than the content of individual settings — has no regulatory precedent, and writing one requires imagining the artifact first. Legislation follows artifacts; the artifact is the paper’s job.

5 The specification

5.1 The defaults manifest

Every shipping application version carries a machine-readable **defaults manifest**: an enumerated list of the user-affecting defaults it installs, each entry with — setting identifier and human-readable name; value; class (privacy, security, data-flow, exposure, cost, behavior); **provenance** (the deciding team or review, the date, the stated rationale — the commit-message discipline for policy); **changed-from** (the previous default, when this one differs); and a **user-facing description** written to the standard the release note aspires to and rarely achieves. The manifest is the lockfile’s civic cousin: what the lockfile is for dependency resolution, the manifest is for choice architecture — the record of what was selected on the user’s behalf, at what version, by whose decision.

5.2 Diff, drift, and the update-time disclosure

The platform diffs manifests across updates and surfaces the delta at update time in the same surface that announces the update — not a settings page the user will never open, but the consent-adjacent moment where the change actually lands: “This update changes 3 defaults: X (was A, now B — reason), Y (...), Z (...).” The disclosure’s grammar is the reconciliation pattern from P-012: a default change without a receiving disclosure is an unrecorded transaction against the user’s configuration estate. **Drift notification** covers the dynamic variant: defaults that change server-side or by remote flag (the feature-flag path that bypasses versioning entirely) — the flag system already knows what it flipped; the ledger makes the flip an event the user’s device can be told about, with a threshold above which silent remote re-defaulting is a violation of the manifest contract itself.

5.3 The personal override registry

The user’s deviations from every product’s defaults become a first-class portable artifact: exportable, versioned, reinstall-surviving, device-portable. The registry is the civilian’s config profile — the MDM insight, pointed at the person instead of the fleet — and it inverts the platform’s strongest friction advantage: your chosen configuration becomes the thing that follows you, and the factory default becomes the thing that must justify itself against it. The registry also enables the setting that does not exist today and should: “**notify me whenever any application changes a default I have deviated from.**”

5.4 Negotiation layers

Defaults compose above the individual: a **household policy** (the parental and shared-device configuration story, intersecting P-013’s delegated operation), an **organizational policy** for the long tail of non-managed devices (the freelancer, the small business, the school), and **importable default packs** — curated, signed, versioned sets of deviations published by parties whose judgment the user has chosen to trust (a privacy NGO’s pack; a photographer’s pack; a low-vision accessibility pack). The pack is the editorial layer defaults never had: a third party whose job is to *represent the user’s interests against the factory configuration* — the consumer-reporting function, instantiated as configuration.

5.5 The default impact assessment

For default changes above an exposure threshold (population affected, behavioral magnitude, class sensitivity), the setter publishes an **impact assessment** — the environmental-impact-statement pattern, applied to choice architecture: the expected behavioral displacement (the A/B machinery already measures this internally with scientific precision; the assessment makes the measurement public), the affected population and its asymmetries (who bears the change’s cost and who its benefit), the reversal path (how a user returns to prior behavior, at what friction cost — making the friction itself an auditable number, per §2’s mechanism 2), and the security/privacy surface delta. The assessment’s honest threshold calibration is a governance question the manifest standard settles per class: a dark-mode default needs none; a telemetry default needs all of it.

5.6 Procurement-grade defaults BoM

The aggregated manifest view for the parties that govern at portfolio scale: the CIO diffing vendors by configuration behavior; the browser or OS store reviewer requiring a manifest as listing condition; the regulator auditing amendment cadence without per-bit legislation; the researcher finally able to run the longitudinal science of default drift that the field has never had data for. The BoM is the artifact that makes the practice of default-setting *studyable* — and per §3’s research gap, the studyability is itself the unlock.

6 Evaluation design

1. **Legibility trials.** Can users parse manifest diffs? Present update-time disclosures across complexity tiers (3 changes vs. 30; plain-language vs. manifest-native); measure comprehension, decision rate, and time. The display contract (P-015’s presentation science) is the dependent discipline; the falsifier is the disclosure collapsing into consent-fatigue noise above some N of changes — which would itself be a finding with regulatory bite.
2. **Drift-detection accuracy.** Instrument a device fleet’s actual default changes over a release season (with consent); measure the manifest-diff’s recall against behavioral detection (settings that observably moved) — the gap between declared and actual amendments is the ecosystem’s honesty measurement, publishable as a score.
3. **Override-registry survival.** Longitudinal trials of registry portability across device migration, OS reinstalls, and app version churn; measure configuration-loss events (the “spend an evening re-doing all my settings after the new phone” ritual, finally counted) and the registry’s kill rate on them.
4. **Default-pack trust dynamics.** Who do users delegate configuration judgment to, how do pack updates propagate consent, and what are the abuse modes (the pack as a new at-

tack surface for preference manipulation — the specification’s own §5.4 power concentration confronted empirically)?

5. **Impact-assessment calibration.** Retro-score published assessments against measured post-change behavior (the A/B records, where disclosed): do the assessments predict displacement honestly, or do they launder? Retroactive audits of famous default changes (ATT 2021; encryption-by-default 2016; Recall 2024’s default-on-to-off reversal under public pressure) against the assessment format — the historical cases the theory must fit.

7 What this is not

This is not a **settings manager** — the cataloguers of settings panels exist; the ledger governs the *amendments*, not the inventory. This is not a **regulation proposal** — the specification runs entirely on a platform volunteering honesty or a market pricing it, though every instrument here is the artifact a regulator would reference the day after the next billion-dollar default change. This is not “**all defaults must be opt-in**” — default-on is often the *public-spirited* choice (encryption; safety updates; spam filtering) and the ledger’s disclosure is precisely what distinguishes legitimate default-setting from laundering; a world where every bit demands a decision is a world where the friction asymmetry (§2.2) simply taxes everyone into permanent non-choice, which no serious defaults literature recommends. This is not an **anti-business instrument** — default power will remain valuable under a ledger; the ledger prices it, bounds it, and moves some of its value to the user, exactly as disclosure regimes moved value from issuers to investors without ending issuance. And this is not a **one-platform campaign** — the manifest is a standard whose power compounds with adoption; the specification’s natural first movers (browsers, OS stores, procurement) each suffice to start the compounding.

8 Objections, confronted

“**Users don’t care; that’s why they don’t change defaults.**” The behavioral literature says the opposite of the objection’s premise: users *behave* on defaults — the organ-donation gap, the enrollment gap, the encryption installed on a billion phones — with effect sizes the persuasion literature can only envy. Not caring is the *measured name* for the friction asymmetry, not evidence of consent; the user who cannot name the bit that moved cannot be surveyed as having chosen it. The ledger does not force users to care; it makes the *amendment* an event the system can count, which is a different thing entirely from demanding a decision per bit.

“**Disclosure at update time is consent fatigue; nobody reads release notes.**” Partly true today — because release notes are engineered to be unread (marketing copy, no diffs, no grammar of change). §6.1 exists precisely to find the disclosure’s actual legibility frontier, and the consent-fatigue literature sets the design constraints (short diffs, thresholds, class-prioritization). The honest comparison is not “disclosure vs. silence” but “N disclosed amendments vs. N silent ones”: the current system delivers the same fatigue *with zero information*, which is the strictly dominated outcome the objection defends.

“**Remote flags and server-side defaults make version manifests obsolete.**” Server-side configuration is exactly why §5.2’s drift notification is a first-class instrument rather than an afterthought: the flag system knows what it flips; the ledger requires the flip to be an *event* with the same provenance grammar as a version change, and the manifest contract treats silent remote re-defaulting above threshold as a violation. The objection identifies the harder half of the problem and hands the specification its most important clause.

“**This adds bureaucracy to design; default-setting will slow to a crawl.**” The impact assessment has a threshold — dark-mode flips, icon changes, and typography never

reach it; the telemetry default and the data-sharing default do, because those are precisely the bits whose silent movement the effect sizes justify governing. The discipline’s cost lands on the changes whose value derives from invisibility; the changes whose value survives disclosure ship as fast as ever. “Slows down laundering” is the accurate summary of the objection, and it is the specification working as intended.

“The manifest will be gamed — false rationale, burying changes in thousands of entries.” The same gaming pressure applies to financial disclosure and nutrition labels, whose regulators answered with standard grammar, spot audits, and penalties for material misstatement — instruments the manifest standard’s enforcement layer inherits as policy, not invention. The research community’s drift-detection work (§6.2) becomes the spot-audit: behavioral detection of undeclared amendments, publishable as ecosystem honesty scores. A gamed ledger is still an attack surface where the current system has *no surface at all* — every known adversarial-disclosure regime outperforms the void it replaced.

“Defaults are the vendor’s product decision; users who dislike them can leave.” For competitive settings, largely true and untouched by the ledger. For the settings that matter — the coordination defaults (§2.4), the exposure and data-flow defaults, the platform-default placements worth billions precisely because leaving does not exit the topology — the exit is fictional and the price is public. The market already prices the default at \$20 billion a year; the ledger is the buyer’s receipt.

9 Limitations

The manifest standard’s coverage will always trail the creativity of configuration surfaces (flags, remote config, ML-personalized quasi-defaults — the system that *learns* your default and quietly implements it, which is a defaults ledger problem wearing a recommendation-system costume, and is noted here as the specification’s known frontier rather than solved). The disclosure’s fatigue frontier (§6.1) is real and will bound the instrument’s civilian uptake; the enterprise/procurement layer is where early adoption lives regardless. The pack mechanism (§5.4) creates a new trust concentration whose abuse modes are genuinely open research (§6.4). The impact assessment inherits the measurement honesty problem of every self-reported disclosure regime, and its calibration (§6.5) depends on A/B disclosure the industry will resist providing. The standard’s political economy is the honest hardest problem: the parties with the machinery are the parties the ledger governs, so adoption arrives through the three channels §4 predicts — regulation, competitive honesty, or research forcing — and the specification’s contribution is to make all three *possible*, not to pick the one that fires. And a defaults ledger does not decide any default; it records who did, why, and what it changed you from. That is a smaller claim than legislation, and it is the necessary claim: the record precedes the argument.

10 Conclusion

Organ donation at 99.98% versus 12% between neighboring countries; retirement saved or not by a checkbox’s prior state; a billion conversations encrypted by a release decision; twenty billion dollars a year for one search box’s prior placement; a trillion-dollar industry repriced by one prompt shipped default-on. No force in the software world moves humans at this scale, and no force of comparable power anywhere operates with less record: no manifest of what was decided, no provenance of who decided, no diff of what changed, no notice that the amendment occurred, no override that survives the next upgrade, no assessment of what the change did to the population it was done to. The enterprise built this governance for itself decades ago because it could price its absence; the public — the polity on whom the defaults are legislated — has been left with vibes and a reset button. This paper specified the civic layer: the manifest, the

diff, the drift event, the portable override, the negotiation and pack layers, the impact statement, the bill of materials. Every fragment is proven; the composition is specified; the first mover — a browser, a store, a platform, a regulator, a procurement office — gets to be the world’s first default polity with a parliamentary record. The defaults were never neutral. Now, at least, there will be a record of that.

References

1. Johnson, E. J., & Goldstein, D. (2003). “Do Defaults Save Lives?” *Science*, 302. [Peer-reviewed; high confidence — the canonical defaults study.]
2. Madrian, B. C., & Shea, D. F. (2001). “The Power of Suggestion: Inertia in 401(k) Participation and Savings Behavior.” *Quarterly Journal of Economics*, 116(4). [Peer-reviewed; high confidence.]
3. WhatsApp end-to-end encryption by default — April 2016, Signal Protocol, ~1B users at flip. [Company/platform record; high confidence.]
4. Apple App Tracking Transparency (2021) and Meta’s published estimate of ~\$10B revenue impact for 2022. [Platform + earnings-call record; high confidence.]
5. Google–Apple default-search payments, ~\$20B for 2022, disclosed in *United States v. Google* exhibits. [Antitrust-trial record; high confidence on the order of magnitude; verification queued for the exact exhibit figure.]
6. Thaler, R. H., & Sunstein, C. R. (2008). *Nudge* — choice architecture and default effects. [Book; high confidence.]
7. Microsoft Windows Recall (2024) — announced default-on, postponed after public reaction, shipped off-by-default. [Press/platform record; high confidence on the sequence.]
8. Group Policy / MDM / managed-preferences lineage — enterprise configuration governance. [Platform/standards records; high confidence.]
9. Lockfile and manifest practice — package-lock, Cargo.lock, and the supply-chain’s diffable default records. [Tooling record; high confidence.]
10. Do Not Track and IE10’s default-on DNT (2012), with the Apache counter-response ignoring IE10’s DNT header. [Standards/infrastructure record; high confidence — the classic default-politics episode.]
11. The browser-ballot remedies in the EU Microsoft case. [Regulatory record; high confidence.]
12. Defaults-in-privacy research line; the endorsement-heuristic literature. [Research lines; high confidence on existence; specific-paper verification queued.]
13. GDPR (2016/679) and U.S. state privacy statutes — per-setting default mandates. [Legislative records; high confidence.]
14. p-rick P-012, *The Complexity Ledger* — the reconciliation pattern (a change without a receiving disclosure is an unrecorded transaction) applied here to configuration estate. [Program cross-reference.]
15. p-rick P-013, *Delegated Operation* — household negotiation layers and the consent grammar for who may amend whose defaults. [Program cross-reference.]
16. p-rick P-015, *The Uncertain Document* — presentation science and the display-contract discipline the disclosure layer borrows. [Program cross-reference.]

Working-paper note: references marked “verification queued” are recorded from domain knowledge at high confidence and await the program’s live citation-verification pass (see agents.md); the verification ledger is maintained with the paper sources.